

Data Intensive Computing - Review Questions 1

Deadline: September 4, 2020

1. **1 point.** Explain how a file region can be left in an inconsistent state in GFS?

2. **1 point.** Briefly explain what are the differences between GFS and DFS?

3. **1 point.** Show with an example that in the CAP theorem, we can have only two properties out of Consistency, Availability, and Partition Tolerance at the same time.

4. **1 point.** How does BigTable provide strong consistency?

5. **1 point.** Write a simple query in Cypher for each of the following requests:

- Match a Person called “John Doe”
- Find FRIENDS_OF “John Doe”
- Count “John Doe”’s direct acquaintances