

# Data Intensive Computing - Review Questions 2

Deadline: September 11, 2020

1. **1 point.** Imagine you are working in a big company, and your company is planning to launch the next big Blogging platform. Tomorrow morning you go to your office and see the following mail from your CEO regarding a new work. How do you answer to this email? Hint: use MapReduce to solve it, and explain what you do in Map and Reduce phases.

Dear <Your Name>,

As you know we are building a blogging platform, and I need some statistics. I need to find out, across all blogs ever written on our blogging system, how many times 1 character words occur (e.g., "a", "I"), How many times two character words occur (e.g., "be", "is"), and so on. I know its a really big job. I am going on a vacation for one week, and its really important that I've this when I return. Good luck.

Regards,

The CEO

- 
2. **1 point.** What is the problem of skewed data in MapReduce? (when the key distribution is skewed)

- 
3. **1 point.** Briefly explain the differences between Map-side join and Reduce-side join in Map-Reduce?

- 
4. **1 point.** Explain briefly why the following code does not work correctly on a cluster of computers. How can we fix it?

```
val uni = sc.parallelize(Seq(("SICS", 1), ("KTH", 2)))
uni.foreach(println)
```

- 
5. **1 point.** Assume you are reading the file `campus.txt` from HDFS with the following format:

```
SICS CSL
KTH CSC
UCL NET
SICS DNA
...
```

Draw the lineage graph for the following code and explain how Spark uses the lineage graph to handle failures.

```
val file = sc.textFile("hdfs://campus.txt")
val pairs = file.map(x => (x.split(" ")(0), x.split(" ")(1)))
val groups = pairs.groupByKey()
val uni = sc.parallelize(Seq(("SICS", 1), ("KTH", 2)))
val joins = groups.join(uni)
val sics = joins.filter(x => x.contains("SICS"))
val list = sics.map(x => x._2)
val result = list.count
```