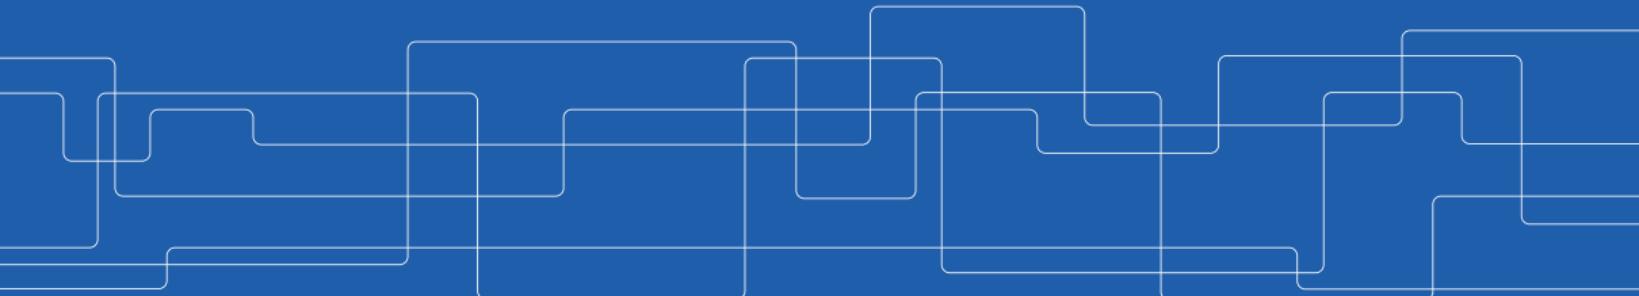




Large Scale Graph Processing - Pregel, GraphLab, and XStream

Amir H. Payberah
payberah@kth.se
08/10/2018






The Course Web Page

<https://id2221kth.github.io>



Where Are We?




- ▶ A **flexible abstraction** for describing relationships between **discrete objects**.





Large Graph





Graph Algorithms Challenges

- ▶ Difficult to extract parallelism based on partitioning of the data.
- ▶ Difficult to express parallelism based on partitioning of computation.




Graph Algorithms Challenges

- ▶ Difficult to extract parallelism based on partitioning of the data.
- ▶ Difficult to express parallelism based on partitioning of computation.
- ▶ Graph partition is a challenging problem.


Graph Partitioning

- ▶ Partition large scale graphs and **distribut to hosts**.




Edge-Cut Graph Partitioning

- ▶ Divide **vertices** of a graph into disjoint clusters.
- ▶ Nearly **equal size** (w.r.t. the number of **vertices**).
- ▶ With the **minimum number of edges** that **span** separated clusters.



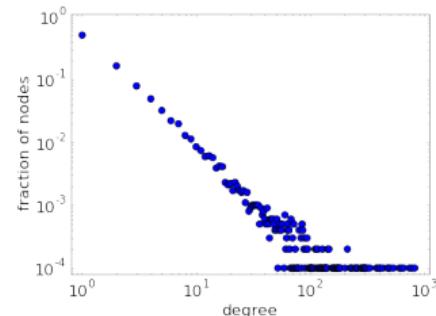
Vertex-Cut Graph Partitioning

- ▶ Divide **edges** of a graph into **disjoint clusters**.
- ▶ Nearly **equal size** (w.r.t. the number of **edges**).
- ▶ With the **minimum** number of **replicated vertices**.




Edge-Cut vs. Vertex-Cut Graph Partitioning (1/2)

- ▶ Natural graphs: skewed **Power-Law** degree distribution.
- ▶ **Edge-cut** algorithms perform **poorly** on Power-Law Graphs.



Edge-Cut vs. Vertex-Cut Graph Partitioning (2/2)





Different Approaches to Process Large Scale Graphs


- ▶ Think like a **vertex**
- ▶ Think like an **edge**
- ▶ Think like a **table**
- ▶ Think like a **graph**
- ▶ Think like a **matrix**



Think Like a Vertex


Think Like a Vertex

- ▶ Each vertex computes **individually** its value (in **parallel**)
- ▶ Computation typically depends on the **neighbors**.





Think Like a Vertex

- ▶ Each vertex computes **individually** its value (in **parallel**)
- ▶ Computation typically depends on the **neighbors**.
- ▶ Also known as **graph-parallel** processing model.




Data-Parallel vs. Graph-Parallel Computation





Pregel

- ▶ Large-scale graph-parallel processing platform developed at Google.
- ▶ Inspired by bulk synchronous parallel (BSP) model.



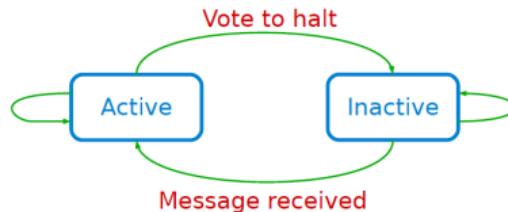


Execution Model (1/2)

- ▶ Applications run in sequence of **iterations**, called **supersteps**
- ▶ A vertex in superstep **S** can:
 - **reads** messages sent to it in superstep **S-1**.
 - **sends** messages to other vertices: receiving at superstep **S+1**.
 - **modifies** its state.
- ▶ Vertices communicate directly with one another by **sending messages**.

Execution Model (2/2)

- ▶ Superstep 0: all vertices are in the **active** state.
- ▶ A vertex **deactivates** itself by voting to **halt**: no further work to do.
- ▶ A halted vertex can be active if it **receives a message**.
- ▶ The whole algorithm terminates when:
 - All vertices are **simultaneously** **inactive**.
 - There are **no messages in transit**.




Example: Max Value (1/4)

```
i_val := val

for each message m
    if m > val then val := m

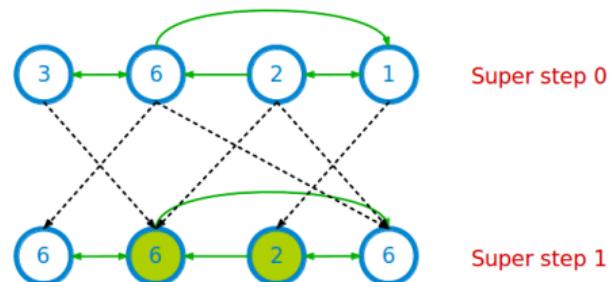
if i_val == val then
    vote_to_halt
else
    for each neighbor v
        send_message(v, val)
```



Super step 0


Example: Max Value (2/4)

```
i_val := val  
  
for each message m  
  if m > val then val := m  
  
if i_val == val then  
  vote_to_halt  
else  
  for each neighbor v  
    send_message(v, val)
```



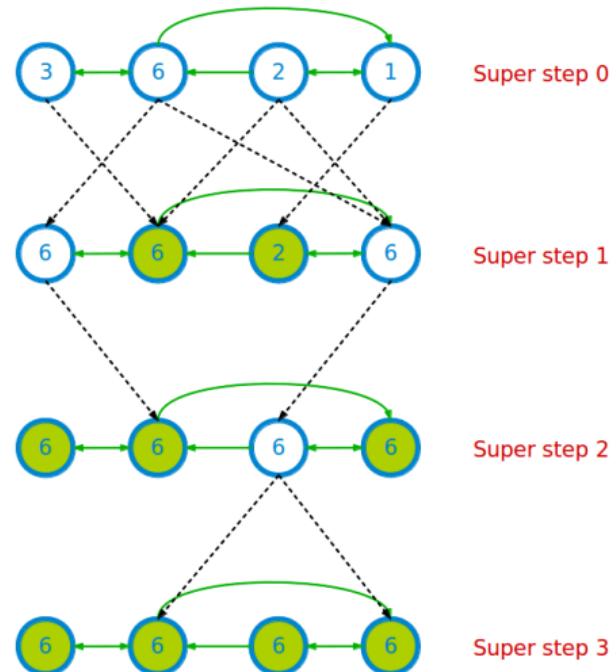
Example: Max Value (3/4)

```
i_val := val  
  
for each message m  
  if m > val then val := m  
  
if i_val == val then  
  vote_to_halt  
else  
  for each neighbor v  
    send_message(v, val)
```

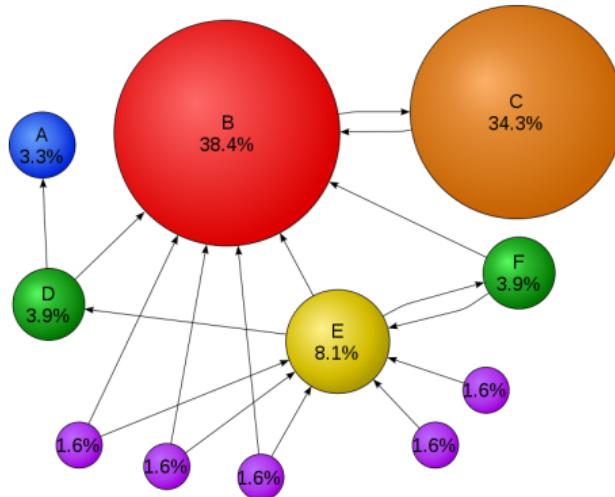


Example: Max Value (4/4)

```
i_val := val  
  
for each message m  
  if m > val then val := m  
  
if i_val == val then  
  vote_to_halt  
else  
  for each neighbor v  
    send_message(v, val)
```



Example: PageRank



$$R[i] = 0.15 + \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$



Example: PageRank

```
Pregel_PageRank(i, messages):
    // receive all the messages
    total = 0
    foreach(msg in messages):
        total = total + msg

    // update the rank of this vertex
    R[i] = 0.15 + total

    // send new messages to neighbors
    foreach(j in out_neighbors[i]):
        sendmsg(R[i] * wij) to vertex j
```

$$R[i] = 0.15 + \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$



Graph Partitioning

- ▶ Edge-cut partitioning
- ▶ The pregel library divides a graph into a number of partitions.
- ▶ Each partition consists of vertices and all of those vertices' outgoing edges.
- ▶ Vertices are assigned to partitions based on their vertex-ID (e.g., $\text{hash}(\text{ID})$).



System Model

- ▶ Master-worker model.
- ▶ The master
 - Coordinates workers.
 - Assigns one or more partitions to each worker.
 - Instructs each worker to perform a superstep.
- ▶ Each worker
 - Executes the local computation method on its vertices.
 - Maintains the state of its partitions.
 - Manages messages to and from other workers.



Fault Tolerance

- ▶ Fault tolerance is achieved through **checkpointing**.
 - Saved to persistent storage
- ▶ At **start of each superstep**, master tells workers to **save** their state:
 - Vertex values, edge values, incoming messages
- ▶ Master saves **aggregator values** (if any).
- ▶ When master **detects** one or more **worker failures**:
 - All workers revert to last **checkpoint**.



Pregel Limitations

- ▶ Inefficient if different regions of the graph converge at **different speed**.
- ▶ Runtime of each phase is determined by the **slowest** machine.




GraphLab/Turi



GraphLab

- ▶ GraphLab allows **asynchronous** iterative computation.
- ▶ **Vertex scope** of vertex v : the data stored in v , and in all **adjacent** vertices and edges.
- ▶ A vertex can **read** and **modify** any of the data in its **scope** (**shared memory**).

- ▶ GraphLab allows **asynchronous** iterative computation.
- ▶ **Vertex scope** of vertex v : the data stored in v , and in all **adjacent vertices and edges**.
- ▶ A vertex can **read** and **modify** any of the data in its **scope** (**shared memory**).





Example: PageRank (GraphLab)

```
GraphLab_PageRank(i)
    // compute sum over neighbors
    total = 0
    foreach(j in in_neighbors(i)):
        total = total + R[j] * wji

    // update the PageRank
    R[i] = 0.15 + total

    // trigger neighbors to run again
    foreach(j in out_neighbors(i)):
        signal vertex-program on j
```

$$R[i] = 0.15 + \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$

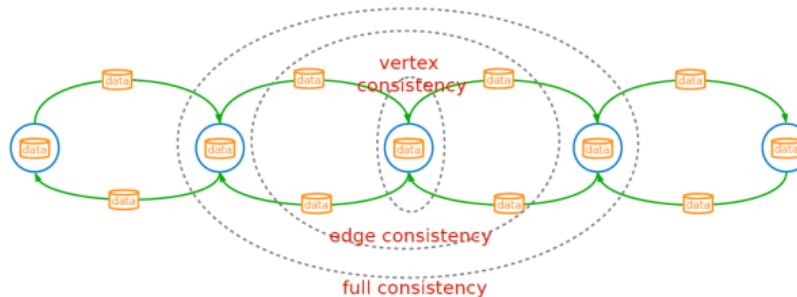


Consistency (1/4)

- ▶ Overlapped scopes: **race-condition** in simultaneous execution of **two update functions**.


Consistency (1/4)

- ▶ Overlapped scopes: **race-condition** in simultaneous execution of **two update functions**.




- ▶ **Full consistency:** during the execution $f(v)$, no other function reads or modifies data within the v scope.

Consistency (2/4)



- ▶ **Edge consistency:** during the execution $f(v)$, no other function reads or modifies any of the data on v or any of the edges adjacent to v .


Consistency (3/4)



- ▶ **Vertex consistency**: during the execution $f(v)$, no other function will be applied to v .

Consistency (4/4)

Consistency vs. Parallelism



[Low, Y., GraphLab: A Distributed Abstraction for Large Scale Machine Learning (Doctoral dissertation, University of California), 2013.]



Consistency Implementation

- ▶ **Distributed locking:** associating a **readers-writer** lock with each vertex.



Consistency Implementation

- ▶ **Distributed locking:** associating a **readers-writer** lock with each vertex.
- ▶ **Vertex consistency**
 - Central vertex (**write-lock**)



Consistency Implementation

- ▶ **Distributed locking:** associating a **readers-writer** lock with each vertex.
- ▶ **Vertex consistency**
 - Central vertex (**write-lock**)
- ▶ **Edge consistency**
 - Central vertex (**write-lock**), Adjacent vertices (**read-locks**)



Consistency Implementation

- ▶ **Distributed locking:** associating a **readers-writer** lock with each vertex.
- ▶ **Vertex consistency**
 - Central vertex (**write-lock**)
- ▶ **Edge consistency**
 - Central vertex (**write-lock**), Adjacent vertices (**read-locks**)
- ▶ **Full consistency**
 - Central vertex (**write-locks**), Adjacent vertices (**write-locks**)




Consistency Implementation

- ▶ **Distributed locking**: associating a **readers-writer** lock with each vertex.
- ▶ **Vertex consistency**
 - Central vertex (**write-lock**)
- ▶ **Edge consistency**
 - Central vertex (**write-lock**), Adjacent vertices (**read-locks**)
- ▶ **Full consistency**
 - Central vertex (**write-locks**), Adjacent vertices (**write-locks**)
- ▶ **Deadlocks** are avoided by acquiring **locks sequentially** following a **canonical order**.

Graph Partitioning

- ▶ Edge-cut partitioning.
- ▶ Two-phase partitioning:
 1. Convert a large graph into a small meta-graph
 2. Partition the meta-graph





Fault Tolerance - Synchronous

- ▶ The systems **periodically** signals all computation activity to **halt**.
- ▶ Then **synchronizes all caches**, and **saves to disk** all data which has been modified since the last snapshot.
- ▶ **Simple**, but eliminates the systems advantage of **asynchronous** computation.



Fault Tolerance - Asynchronous

- ▶ Based on the Chandy-Lamport algorithm.
- ▶ The **snapshot** function is implemented **as a function in vertices**.
 - It takes **priority** over all other update functions.

```
if v was already snapshotted then
    ↘ Quit
    Save  $D_v$  // Save current vertex
    // Save all edges connected to un-snapshotted vertices
foreach  $u \in N[v]$  do                                // Loop over neighbors
    if u was not snapshotted then
        ↘ Save  $D_{u \rightarrow v}$  if edge  $u \rightarrow v$  exists
        ↘ Save  $D_{v \rightarrow u}$  if edge  $v \rightarrow u$  exists
        ↘ Reschedule u for a Snapshot Update
    Mark v as snapshotted
```



GraphLab2/Turi (PowerGraph)



PowerGraph

- ▶ Factorizes the **local vertices** functions into the **Gather**, **Apply** and **Scatter** phases.
- ▶ **Vertx-cut** partitioning.



Programming Model

- ▶ Gather-Apply-Scatter (GAS)
- ▶ **Gather**: accumulate information from neighborhood.
- ▶ **Apply**: apply the accumulated value to center vertex.
- ▶ **Scatter**: update adjacent edges and vertices.



Execution Model (1/2)

- ▶ Initially **all vertices** are **active**.
- ▶ It executes the **vertex-program** on the **active vertices** until none remain.
 - Once a vertex-program completes the **scatter** phase it becomes **inactive** until it is reactivated.
 - Vertices can activate **themselves** and **neighboring vertices**.



Execution Model (1/2)

- ▶ Initially **all vertices** are **active**.
- ▶ It executes the **vertex-program** on the **active vertices** until none remain.
 - Once a vertex-program completes the **scatter** phase it becomes **inactive** until it is reactivated.
 - Vertices can activate **themselves** and **neighboring vertices**.
- ▶ PowerGraph can execute both **synchronously** and **asynchronously**.



Execution Model (2/2)

- ▶ **Synchronous** scheduling like Pregel.
 - Executing the **gather**, **apply**, and **scatter** **in order**.
 - Changes made to the vertex/edge data are committed at the **end** of each step.



Execution Model (2/2)

► **Synchronous** scheduling like **Pregel**.

- Executing the **gather**, **apply**, and **scatter** in order.
- Changes made to the vertex/edge data are committed at the **end** of each step.

► **Asynchronous** scheduling like **GraphLab**.

- Changes made to the vertex/edge data during the **apply** and **scatter** functions are **immediately** committed to the graph.
- **Visible** to subsequent computation on neighboring vertices.



Example: PageRank (PowerGraph)

```
PowerGraph_PageRank(i):
    Gather(j -> i):
        return wji * R[j]

    sum(a, b):
        return a + b

    // total: Gather and sum
    Apply(i, total):
        R[i] = 0.15 + total

    Scatter(i -> j):
        if R[i] changed then activate(j)
```

$$R[i] = 0.15 + \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$



Graph Partitioning (1/2)

- ▶ **Random** vertex-cuts
- ▶ Randomly assign edges to machines.
- ▶ Completely parallel and easy to **distribute**.
- ▶ High replication factor.



Graph Partitioning (2/2)

- ▶ Greedy vertex-cuts
- ▶ $A(v)$: set of machines that vertex v spans.

Graph Partitioning (2/2)

- ▶ Greedy vertex-cuts
- ▶ $A(v)$: set of machines that vertex v spans.
- ▶ Case 1: If $A(u) \cap A(v) \neq \emptyset$, then the edge (u, v) should be assigned to a machine in the intersection.



Graph Partitioning (2/2)

- ▶ Greedy vertex-cuts
- ▶ $A(v)$: set of machines that vertex v spans.
- ▶ Case 1: If $A(u) \cap A(v) \neq \emptyset$, then the edge (u, v) should be assigned to a machine in the intersection.
- ▶ Case 2: If $A(u) \cap A(v) = \emptyset$, then the edge (u, v) should be assigned to one of the machines from the vertex with the most unassigned edges.

Graph Partitioning (2/2)

- ▶ Greedy vertex-cuts
- ▶ $A(v)$: set of machines that vertex v spans.
- ▶ Case 1: If $A(u) \cap A(v) \neq \emptyset$, then the edge (u, v) should be assigned to a machine in the intersection.
- ▶ Case 2: If $A(u) \cap A(v) = \emptyset$, then the edge (u, v) should be assigned to one of the machines from the vertex with the most unassigned edges.
- ▶ Case 3: If only one of the two vertices has been assigned, then choose a machine from the assigned vertex.

Graph Partitioning (2/2)

- ▶ Greedy vertex-cuts
- ▶ $A(v)$: set of machines that vertex v spans.
- ▶ Case 1: If $A(u) \cap A(v) \neq \emptyset$, then the edge (u, v) should be assigned to a machine in the intersection.
- ▶ Case 2: If $A(u) \cap A(v) = \emptyset$, then the edge (u, v) should be assigned to one of the machines from the vertex with the most unassigned edges.
- ▶ Case 3: If only one of the two vertices has been assigned, then choose a machine from the assigned vertex.
- ▶ Case 4: If $A(u) = A(v) = \emptyset$, then assign the edge (u, v) to the least loaded machine.



Think Like an Edge

Motivation

- ▶ Could we process **massive graphs** on a **single machine**?
- ▶ **Disk-based** processing
 - Graph traversal = **random access**
 - Random access is **inefficient** for storage

Medium	Read (MB/s)		Write (MB/s)	
	Random	Sequential	Random	Sequential
RAM	567	2605	1057	2248
SSD	22.64	355	49.16	298
Disk	0.61	174	1.27	170

Note: 64 byte cachelines, 4K blocks (disk random), 16M chunks
(disk sequential)

Eiko Y., and Roy A., "Scale-up Graph Processing: A Storage-centric View", 2013.



X-Stream




X-Stream

- ▶ X-Stream makes graph accesses **sequential**.
- ▶ Contribution:
 - Edge-centric scatter-gather model
 - Streaming partitions

Vertex-Centric vs. Edge-Centric Programming Model (1/2)

- ▶ Vertex-centric gather-scatter: iterates over vertices
- ▶ Edge-centric gather-scatter: iterates over edges





Vertex-Centric vs. Edge-Centric Programming Model (2/2)

```
Until convergence {
    // the scatter phase
    for all vertices v that need to scatter updates
        send updates over outgoing edges of v

    // the gather phase
    for all vertices v that have updates
        apply updates from inbound edges of v
}
```



Vertex-Centric vs. Edge-Centric Programming Model (2/2)


```
Until convergence {
    // the scatter phase
    for all vertices v that need to scatter updates
        send updates over outgoing edges of v

    // the gather phase
    for all vertices v that have updates
        apply updates from inbound edges of v
}
```

```
Until convergence {
    // the scatter phase
    for all edges e
        send update over e

    // the gather phase
    for all edges e that have updates
        apply update to e.destination
}
```

Vertex-Centric Breadth First Search (1/5)



edges	
src	dest
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
8	5
8	6


vertices

v
1
2
3
4
5
6
7
8

```
Until convergence {
    // the scatter phase
    for all vertices v that need to scatter updates
        send updates over outgoing edges of v

    // the gather phase
    for all vertices v that have updates
        apply updates from inbound edges of v
}
```

Vertex-Centric Breadth First Search (2/5)



vertices

v
1
2
3
4
5
6
7
8


edges

src	dest
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
8	5
8	6

```
Until convergence {
    // the scatter phase
    for all vertices v that need to scatter updates
        send updates over outgoing edges of v

    // the gather phase
    for all vertices v that have updates
        apply updates from inbound edges of v
}
```

Vertex-Centric Breadth First Search (3/5)



edges	
src	dest
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
8	5
8	6

vertices

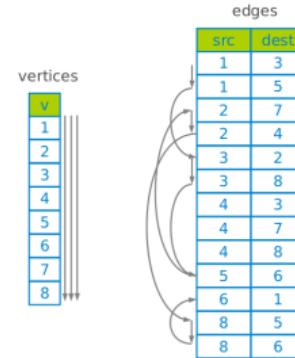

v
1
2
3
4
5
6
7
8

↓↓↓

```
Until convergence {
    // the scatter phase
    for all vertices v that need to scatter updates
        send updates over outgoing edges of v

    // the gather phase
    for all vertices v that have updates
        apply updates from inbound edges of v
}
```



Vertex-Centric Breadth First Search (4/5)



```
Until convergence {
    // the scatter phase
    for all vertices v that need to scatter updates
        send updates over outgoing edges of v

    // the gather phase
    for all vertices v that have updates
        apply updates from inbound edges of v
}
```


Vertex-Centric Breadth First Search (5/5)



```
Until convergence {
    // the scatter phase
    for all vertices v that need to scatter updates
        send updates over outgoing edges of v

    // the gather phase
    for all vertices v that have updates
        apply updates from inbound edges of v
}
```

Edge-Centric Breadth First Search (1/5)




v
1
2
3
4
5
6
7
8

edges	
src	dest
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
8	5
8	6

```
Until convergence {
    // the scatter phase
    for all edges e
        send update over e

    // the gather phase
    for all edges e that have updates
        apply update to e.destination
}
```

Edge-Centric Breadth First Search (2/5)



edges	
src	dest
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
8	5
8	6


vertices

v
1
2
3
4
5
6
7
8

edges

```
Until convergence {  
    // the scatter phase  
    for all edges e  
        send update over e  
  
    // the gather phase  
    for all edges e that have updates  
        apply update to e.destination  
}
```

Edge-Centric Breadth First Search (3/5)



vertices

v	1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---	---

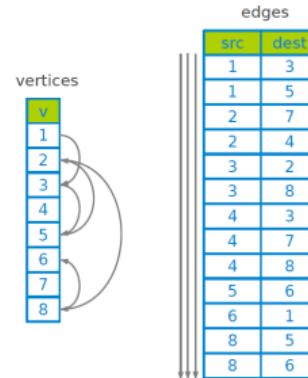
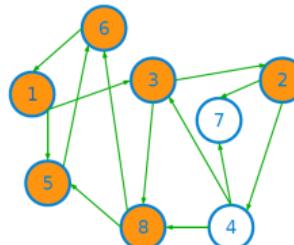
edges

src	dest
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
8	5
8	6

```
Until convergence {
    // the scatter phase
    for all edges e
        send update over e

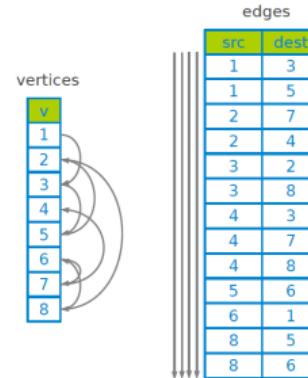

    // the gather phase
    for all edges e that have updates
        apply update to e.destination
}
```

Edge-Centric Breadth First Search (4/5)



```
Until convergence {  
    // the scatter phase  
    for all edges e  
        send update over e  
  
    // the gather phase  
    for all edges e that have updates  
        apply update to e.destination  
}
```

Edge-Centric Breadth First Search (5/5)



```
Until convergence {  
    // the scatter phase  
    for all edges e  
        send update over e  
  
    // the gather phase  
    for all edges e that have updates  
        apply update to e.destination  
}
```




Vertex-Centric vs. Edge-Centric Tradeoff

- ▶ Vertex-centric scatter-gather: $\frac{\text{EdgeData}}{\text{RandomAccessBandwidth}}$
- ▶ Edge-centric scatter-gather: $\frac{\text{Scatters} \times \text{EdgeData}}{\text{SequentialAccessBandwidth}}$
- ▶ Sequential Access Bandwidth \gg Random Access Bandwidth.
- ▶ Few scatter gather iterations for real world graphs.

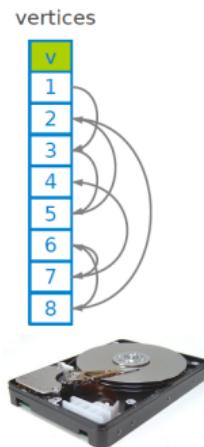
Streaming Partitions (1/4)

- ▶ **Problem:** still have **random** access to **vertex set**.



Streaming Partitions (1/4)




- ▶ **Problem:** still have **random** access to **vertex set**.



Solution

Partition the graph into **streaming partitions**.

Streaming Partitions (2/4)





Streaming Partitions (3/4)

- ▶ A **streaming partition** consists of: a [vertex set](#), an [edge list](#), and an [update list](#).



Streaming Partitions (3/4)

- ▶ A **streaming partition** consists of: a **vertex set**, an **edge list**, and an **update list**.
- ▶ The **vertex set**: a **subset of the vertex set** of the graph that fits into the **memory**.
 - Vertex sets are **mutually disjoint**.
 - Their **union** equals the vertex set of the **entire graph**.



Streaming Partitions (3/4)

- ▶ A **streaming partition** consists of: a **vertex set**, an **edge list**, and an **update list**.
- ▶ The **vertex set**: a **subset of the vertex set** of the graph that fits into the **memory**.
 - Vertex sets are **mutually disjoint**.
 - Their **union** equals the vertex set of the **entire graph**.
- ▶ The **edge list**: all edges whose **source vertex** is in the **partition's vertex set**.



Streaming Partitions (3/4)

- ▶ A **streaming partition** consists of: a **vertex set**, an **edge list**, and an **update list**.
- ▶ The **vertex set**: a **subset of the vertex set** of the graph that fits into the **memory**.
 - Vertex sets are **mutually disjoint**.
 - Their **union** equals the vertex set of the **entire graph**.
- ▶ The **edge list**: all edges whose **source vertex** is in the **partition's vertex set**.
- ▶ The **update list**: all updates whose **destination vertex** is in the **partition's vertex set**.

Streaming Partitions (4/4)

```
// Scatter phase:  
for each streaming_partition p  
    read in vertex set of p  
    for each edge e in edge list of p  
        append update to Uout  
  
// shuffle phase:  
for each update u in Uout  
    let p = partition containing target of u  
    append u to Uin(p)  
destroy Uout  
  
//gather phase:  
for each streaming_partition p  
    read in vertex set of p  
    for each update u in Uin(p)  
        edge_gather(u)  
    destroy Uin(p)
```



Summary



Summary

- ▶ Think like a vertex
 - Pregel: BSP, synchronous parallel model, message passing, edge-cut
 - GraphLab: asynchronous model, shared memory, edge-cut
 - PowerGraph: synchronous/asynchronous model, GAS, vertex-cut
- ▶ Think like an edge
 - XStream: edge-centric GAS, streaming partition



References

- ▶ G. Malewicz et al., “Pregel: a system for large-scale graph processing”, ACM SIGMOD 2010
- ▶ Y. Low et al., “Distributed GraphLab: a framework for machine learning and data mining in the cloud”, VLDB 2012
- ▶ J. Gonzalez et al., “Powergraph: distributed graph-parallel computation on natural graphs”, OSDI 2012
- ▶ A. Roy et al., “X-stream: Edge-centric graph processing using streaming partitions”, ACM SOSP 2013.



Questions?