



# Large Scale Graph Processing - Pregel and GraphLab

Amir H. Payberah  
payberah@kth.se  
01/10/2018

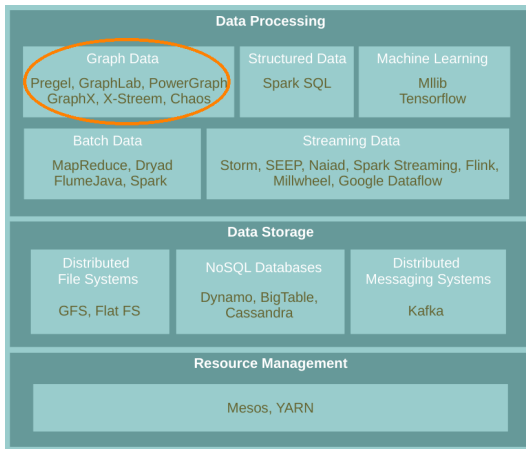




## The Course Web Page

<https://id2221kth.github.io>


# Where Are We?



- ▶ A flexible abstraction for describing relationships between discrete objects.



# Large Graph






# Graph Algorithms Challenges

- ▶ Difficult to extract parallelism based on partitioning of the data.
- ▶ Difficult to express parallelism based on partitioning of computation.
- ▶ Graph partition is a challenging problem.


# Graph Partitioning

- Partition large scale graphs and distribute to hosts.



# Edge-Cut Graph Partitioning


- ▶ Divide **vertices** of a graph into **disjoint clusters**.
- ▶ Nearly **equal size** (w.r.t. the number of **vertices**).
- ▶ With the **minimum number of edges** that **span separated clusters**.






# Vertex-Cut Graph Partitioning


- ▶ Divide **edges** of a graph into **disjoint clusters**.
- ▶ Nearly **equal size** (w.r.t. the number of **edges**).
- ▶ With the **minimum** number of **replicated vertices**.



-




# PageRank with MapReduce



$$R[i] = \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$

► Input




V4: [0.25, V1, V3]

- ▶ Share the rank among all outgoing links

$$V4: (V1, 0.25/2), (V3, 0.25/2)$$

## PageRank Example (2/2)

$$\triangleright R[i] = \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$



V1: (V2, 0.25/3), (V3, 0.25/3), (V4, 0.25/3)

V2: (V3, 0.25/2), (V4, 0.25/2)

V3: (V1, 0.25/1)

V4: (V1, 0.25/2), (V3, 0.25/2)

**► Output after one iteration**

V1: [0.37, V2, V3, V4]

V2: [0.08, V3, V4]

V3: [0.33, V1]

V4: [0.20, V1, V3]

# PageRank in MapReduce - Map (1/2)


## ► Map function

```
map(key: [url, pagerank], value: outlink_list)
  for each outlink in outlink_list:
    emit(key: outlink, value: pagerank / size(outlink_list))

emit(key: url, value: outlink_list)
```

## ► Input (key, value)


```
((V1, 0.25), [V2, V3, V4])
((V2, 0.25), [V3, V4])
((V3, 0.25), [V1])
((V4, 0.25), [V1, V3])
```





## PageRank in MapReduce - Map (2/2)

### ► Map function



```
map(key: [url, pagerank], value: outlink_list)
  for each outlink in outlink_list:
    emit(key: outlink, value: pagerank / size(outlink_list))

emit(key: url, value: outlink_list)
```

### ► Intermediate (key, value)

```
(V2, 0.25/3), (V3, 0.25/3), (V4, 0.25/3), (V3, 0.25/2), (V4, 0.25/2), (V1, 0.25/1),
(V1, 0.25/2), (V3, 0.25/2)
(V1, [V2, V3, V4])
(V2, [V3, V4])
(V3, [V1])
(V4, [V1, V3])
```

## PageRank in MapReduce - Shuffle

### ► Intermediate (key, value)

```
(V2, 0.25/3), (V3, 0.25/3), (V4, 0.25/3), (V3, 0.25/2), (V4, 0.25/2), (V1, 0.25/1),  
(V1, 0.25/2), (V3, 0.25/2)  
(V1, [V2, V3, V4])  
(V2, [V3, V4])  
(V3, [V1])  
(V4, [V1, V3])
```

### ► After shuffling

```
(V1, 0.25/1), (V1, 0.25/2), (V1, [V2, V3, V4])  
(V2, 0.25/3), (V2, [V3, V4])  
(V3, 0.25/3), (V3, 0.25/2), (V3, 0.25/2), (V3, [V1])  
(V4, 0.25/3), (V4, 0.25/2), (V4, [V1, V3])
```

## PageRank in MapReduce - Reduce (1/2)

### ► Reduce function

```
reducer(key: url, value: list_pr_or_urls)
    outlink_list = []
    pagerank = 0

    for each pr_or_urls in list_pr_or_urls:
        if is_list(pr_or_urls):
            outlink_list = pr_or_urls
        else
            pagerank += pr_or_urls

    emit(key: [url, pagerank], value: outlink_list)
```

### ► Input of the Reduce function

```
(V1, 0.25/1), (V1, 0.25/2), (V1, [V2, V3, V4])
(V2, 0.25/3), (V2, [V3, V4])
(V3, 0.25/3), (V3, 0.25/2), (V3, 0.25/2), (V3, [V1])
(V4, 0.25/3), (V4, 0.25/2), (V4, [V1, V3])
```

## PageRank in MapReduce - Reduce (2/2)

### ► Reduce function

```
reducer(key: url, value: list_pr_or_urls)
    outlink_list = []
    pagerank = 0

    for each pr_or_urls in list_pr_or_urls:
        if is_list(pr_or_urls):
            outlink_list = pr_or_urls
        else
            pagerank += pr_or_urls

    emit(key: [url, pagerank], value: outlink_list)
```

### ► Output

```
((V1, 0.37), [V2, V3, V4])
((V2, 0.08), [V3, V4])
((V3, 0.33), [V1])
((V4, 0.20), [V1, V3])
```




# Problems with MapReduce for Graph Analytics

- ▶ MapReduce does **not directly support iterative** algorithms.
  - Invariant graph-topology-data **re-loaded** and **re-processed** at each iteration is **wasting** I/O, network bandwidth, and CPU
- ▶ **Materializations** of intermediate results at every MapReduce iteration **harm performance**.


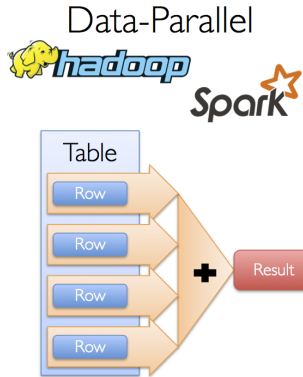
# Think Like a Vertex

# Think Like a Vertex

- ▶ Each vertex computes **individually** its value (in **parallel**).
- ▶ Computation typically depends on the **neighbors**.
- ▶ Also known as **graph-parallel** processing model.




# Data-Parallel vs. Graph-Parallel Computation





# Pregel

- ▶ Large-scale **graph-parallel** processing platform developed at Google.
- ▶ Inspired by **bulk synchronous parallel (BSP)** model.





## Execution Model (1/2)


- ▶ Applications run in sequence of **iterations**, called **supersteps**.
- ▶ A vertex in superstep **S** can:
  - **reads** messages sent to it in superstep **S-1**.
  - **sends** messages to other vertices: receiving at superstep **S+1**.
  - **modifies** its state.
- ▶ Vertices communicate directly with one another by **sending messages**.



- 
- ```
graph LR; Active[Active] -- "Vote to halt" --> Inactive[Inactive]; Inactive -- "Message received" --> Active; Active -- "Message received" --> Active; Inactive -- "Vote to halt" --> Inactive;
```

## Example: Max Value (1/4)

```
i_val := val  
  
for each message m  
  if m > val then val := m  
  
if i_val == val then  
  vote_to_halt  
else  
  for each neighbor v  
    send_message(v, val)
```



Super step 0

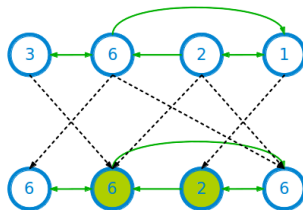
## Example: Max Value (2/4)

```

i_val := val

for each message m
  if m > val then val := m

if i_val == val then
  vote_to_halt
else
  for each neighbor v
    send_message(v, val)
  
```



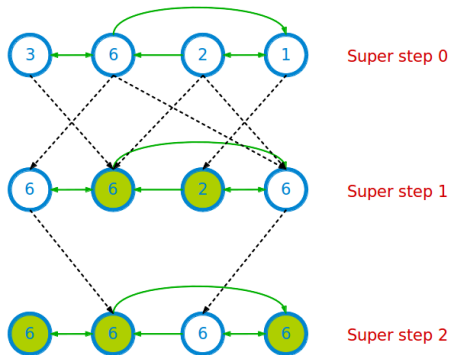
## Example: Max Value (3/4)

```

i_val := val

for each message m
  if m > val then val := m

if i_val == val then
  vote_to_halt
else
  for each neighbor v
    send_message(v, val)
  
```



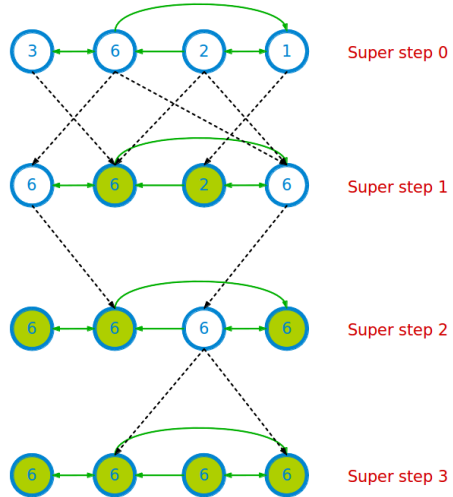
## Example: Max Value (4/4)

```

i_val := val

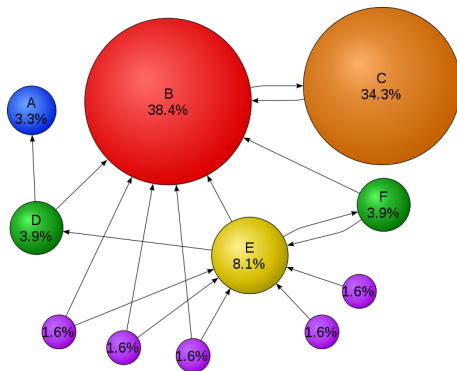
for each message m
  if m > val then val := m

if i_val == val then
  vote_to_halt
else
  for each neighbor v
    send_message(v, val)
  
```





## Example: PageRank



$$R[i] = \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$

## Example: PageRank

```
Pregel_PageRank(i, messages):  
    // receive all the messages  
    total = 0  
    foreach(msg in messages):  
        total = total + msg  
  
    // update the rank of this vertex  
    R[i] = total  
  
    // send new messages to neighbors  
    foreach(j in out_neighbors[i]):  
        sendmsg(R[i] * wij) to vertex j
```

$$R[i] = \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$



# Graph Partitioning

- ▶ Edge-cut partitioning
- ▶ The pregel library divides a graph into a number of partitions.
- ▶ Each partition consists of vertices and all of those vertices' outgoing edges.
- ▶ Vertices are assigned to partitions based on their vertex-ID (e.g.,  $\text{hash}(\text{ID})$ ).



# System Model

- ▶ Master-worker model.
- ▶ The master
  - Coordinates workers.
  - Assigns one or more partitions to each worker.
  - Instructs each worker to perform a superstep.
- ▶ Each worker
  - Executes the local computation method on its vertices.
  - Maintains the state of its partitions.
  - Manages messages to and from other workers.

# Fault Tolerance

- ▶ Fault tolerance is achieved through **checkpointing**.
  - Saved to persistent storage
- ▶ At **start of each superstep**, master tells workers to **save** their state:
  - Vertex values, edge values, incoming messages
- ▶ Master saves **aggregator values** (if any).
- ▶ When master **detects** one or more **worker failures**:
  - All workers revert to last **checkpoint**.

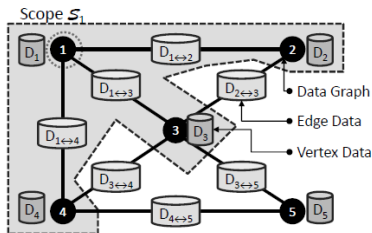


## Pregel Limitations

- ▶ **Inefficient** if different regions of the graph converge at **different speed**.
- ▶ Runtime of each phase is determined by the **slowest** machine.

# GraphLab/Turi

- ▶ GraphLab allows **asynchronous** iterative computation.
- ▶ **Vertex scope** of **vertex  $v$** : the data stored in  $v$ , and in all **adjacent vertices and edges**.
- ▶ A vertex can **read** and **modify** any of the data in its **scope** (**shared memory**).





## Example: PageRank (GraphLab)

```
GraphLab_PageRank(i)
    // compute sum over neighbors
    total = 0
    foreach(j in in_neighbors(i)):
        total = total + R[j] * wji


    // update the PageRank
    R[i] = total

    // trigger neighbors to run again
    foreach(j in out_neighbors(i)):
        signal vertex-program on j
```


$$R[i] = \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$

# Consistency (1/5)

- **Overlapped scopes:** **race-condition** in simultaneous execution of **two update functions**.




## Consistency (2/5)



- **Full consistency:** during the execution  $f(v)$ , no other function reads or modifies data within the  $v$  scope.

## Consistency (3/5)




- **Edge consistency:** during the execution  $f(v)$ , no other function reads or modifies any of the data on  $v$  or any of the edges adjacent to  $v$ .



- **Vertex consistency:** during the execution  $f(v)$ , no other function will be applied to  $v$ .

# Consistency (5/5)

## Consistency vs. Parallelism



[Low, Y., GraphLab: A Distributed Abstraction for Large Scale Machine Learning (Doctoral dissertation, University of California), 2013.]

# Consistency Implementation

- ▶ **Distributed locking**: associating a **readers-writer** lock with each vertex.
- ▶ **Vertex consistency**
  - Central vertex (**write-lock**)
- ▶ **Edge consistency**
  - Central vertex (**write-lock**), Adjacent vertices (**read-locks**)
- ▶ **Full consistency**
  - Central vertex (**write-locks**), Adjacent vertices (**write-locks**)
- ▶ **Deadlocks** are avoided by acquiring **locks sequentially** following a **canonical order**.

- 
- The diagram illustrates the construction of a meta-graph from a graph. The graph is partitioned into three clusters (blue ovals). The meta-graph has nodes representing these clusters, with edges labeled by the weights of the edges between nodes in the original graph.





## Fault Tolerance - Synchronous

- ▶ The systems **periodically** signals all computation activity to **halt**.
- ▶ Then **synchronizes all caches**, and **saves to disk** all data which has been modified since the last snapshot.
- ▶ **Simple**, but eliminates the systems advantage of **asynchronous** computation.

# Fault Tolerance - Asynchronous

- ▶ Based on the **Chandy-Lamport** algorithm.
- ▶ The **snapshot** function is implemented **as a function in vertices**.
  - It takes **priority** over all other update functions.

```
if v was already snapshotted then  
  | Quit  
Save  $D_v$  // Save current vertex  
// Save all edges connected to un-snapshotted vertices  
foreach  $u \in N[v]$  do // Loop over neighbors  
  | if u was not snapshotted then  
    | Save  $D_{u \rightarrow v}$  if edge  $u \rightarrow v$  exists  
    | Save  $D_{v \rightarrow u}$  if edge  $v \rightarrow u$  exists  
    | Reschedule u for a Snapshot Update  
Mark v as snapshotted
```

# GraphLab2/Turi (PowerGraph)



# PowerGraph

- Factorizes the local vertices functions into the Gather, Apply and Scatter phases.



# Programming Model

- ▶ **Gather-Apply-Scatter (GAS)**
- ▶ **Gather:** **accumulate** information from neighborhood.
- ▶ **Apply:** **apply** the accumulated value to center vertex.
- ▶ **Scatter:** **update** adjacent edges and vertices.



## Execution Model (1/2)

- ▶ Initially **all vertices** are **active**.
- ▶ It executes the **vertex-program** on the **active vertices** until none remain.
  - Once a vertex-program completes the **scatter** phase it becomes **inactive** until it is reactivated.
  - Vertices can activate **themselves** and **neighboring** vertices.
- ▶ PowerGraph can execute both **synchronously** and **asynchronously**.

## Execution Model (2/2)

- ▶ **Synchronous** scheduling like **Pregel**.
  - Executing the **gather**, **apply**, and **scatter** in order.
  - Changes made to the vertex/edge data are committed at the **end** of each step.
  
- ▶ **Asynchronous** scheduling like **GraphLab**.
  - Changes made to the vertex/edge data during the **apply** and **scatter** functions are **immediately** committed to the graph.
  - **Visible** to subsequent computation on neighboring vertices.

## Example: PageRank (PowerGraph)

```
PowerGraph_PageRank(i):  
  Gather(j -> i):  
    return wji * R[j]  
  
  sum(a, b):  
    return a + b  
  
  // total: Gather and sum  
  Apply(i, total):  
    R[i] = total  
  
  Scatter(i -> j):  
    if R[i] changed then activate(j)
```

$$R[i] = \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$





## Graph Partitioning (1/2)

- ▶ Vertex-cut partitioning.
- ▶ Random vertex-cuts: randomly assign edges to machines.
- ▶ Completely parallel and easy to distribute.
- ▶ High replication factor.

## Graph Partitioning (2/2)

- ▶ Greedy vertex-cuts
- ▶  $A(v)$ : set of machines that vertex  $v$  spans.
- ▶ Case 1: If  $A(u) \cap A(v) \neq \emptyset$ , then the edge  $(u, v)$  should be assigned to a machine in the intersection.
- ▶ Case 2: If  $A(u) \cap A(v) = \emptyset$ , then the edge  $(u, v)$  should be assigned to one of the machines from the vertex with the most unassigned edges.
- ▶ Case 3: If only one of the two vertices has been assigned, then choose a machine from the assigned vertex.
- ▶ Case 4: If  $A(u) = A(v) = \emptyset$ , then assign the edge  $(u, v)$  to the least loaded machine.

# Summary



# Summary

- ▶ Think like a vertex
  - Pregel: BSP, synchronous parallel model, message passing, edge-cut
  - GraphLab: asynchronous model, shared memory, edge-cut
  - PowerGraph: synchronous/asynchronous model, GAS, vertex-cut

- ▶ G. Malewicz et al., “Pregel: a system for large-scale graph processing”, ACM SIGMOD 2010
- ▶ Y. Low et al., “Distributed GraphLab: a framework for machine learning and data mining in the cloud”, VLDB 2012
- ▶ J. Gonzalez et al., “Powergraph: distributed graph-parallel computation on natural graphs”, OSDI 2012

Questions?