



# Resource Management - Mesos and YARN

Amir H. Payberah  
payberah@kth.se  
2020-09-30



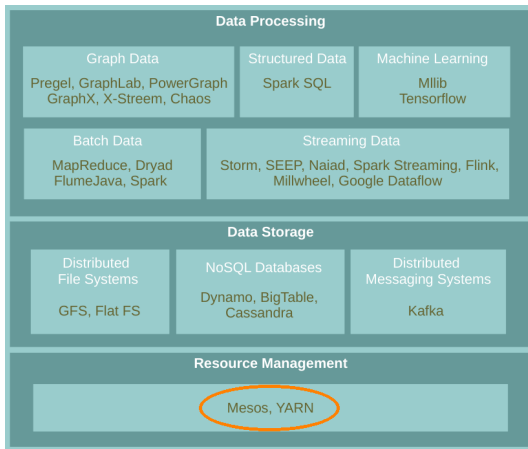


## The Course Web Page

<https://id2221kth.github.io>

<https://tinyurl.com/y4qph82u>

# Where Are We?





# Motivation

- ▶ Rapid innovation in cloud computing.
- ▶ No single framework optimal for all applications.
- ▶ Running each framework on its dedicated cluster:
  - Expensive
  - Hard to share data



## Proposed Solution

- ▶ Running **multiple frameworks** on a **single cluster**.
- ▶ Maximize **utilization** and **share** data between frameworks.
- ▶ Two resource management systems:
  - Mesos
  - YARN


# Mesos

- **Mesos** is a common **resource sharing** layer, over which diverse frameworks can run.



# Computation Model

- ▶ A **framework** (e.g., Hadoop, Spark) manages and runs one or more **jobs**.
- ▶ A **job** consists of one or more **tasks**.
- ▶ A **task** (e.g., map, reduce) consists of one or more **processes** running on same machine.








# Mesos Design Elements

- ▶ Fine-grained sharing
- ▶ Resource offers


# Fine-Grained Sharing

- Allocation at the level of **tasks** within a **job**.
- Improves utilization, latency, and data locality.



# Resource Offer

- Offer available resources to frameworks, let them pick which resources to use and which tasks to launch.
- Keeps Mesos simple, lets it support future frameworks.



## Question?

How to **schedule** resource offering among **frameworks**?



# Schedule Frameworks

- ▶ Global scheduler
- ▶ Distributed scheduler

# Global Scheduler (1/2)

## ► Job requirements


- Response time
- Throughput
- Availability

## ► Job execution plan

- Task DAG
- Inputs/outputs

## ► Estimates

- Task duration
- Input sizes
- Transfer sizes





## Global Scheduler (2/2)


### ► Advantages

- Can achieve **optimal** schedule.

### ► Disadvantages

- **Complexity**: hard to scale and ensure resilience.
- Hard to anticipate **future frameworks** requirements.
- Need to **refactor** existing frameworks.

## Distributed Scheduler (1/3)







## Distributed Scheduler (2/3)

- ▶ Master sends resource offers to frameworks.
- ▶ Frameworks select which offers to accept and which tasks to run.
- ▶ Unit of allocation: resource offer
  - Vector of available resources on a node
  - For example, node1:  $\langle 1\text{CPU}, 1\text{GB} \rangle$ , node2:  $\langle 4\text{CPU}, 16\text{GB} \rangle$



## Distributed Scheduler (3/3)

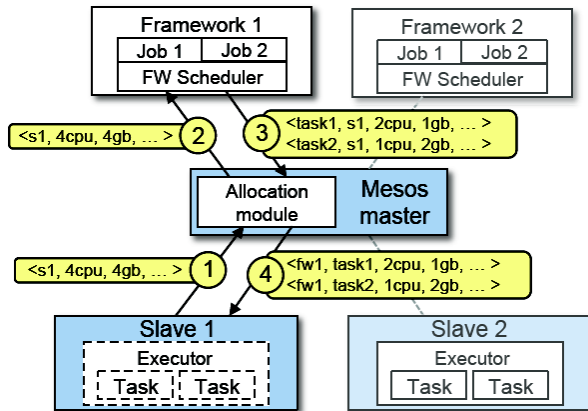
### ► Advantages

- **Simple**: easier to scale and make resilient.
- **Easy to port** existing frameworks, support new ones.

### ► Disadvantages


- Distributed scheduling decision: **not optimal**.

# Mesos Architecture (1/4)



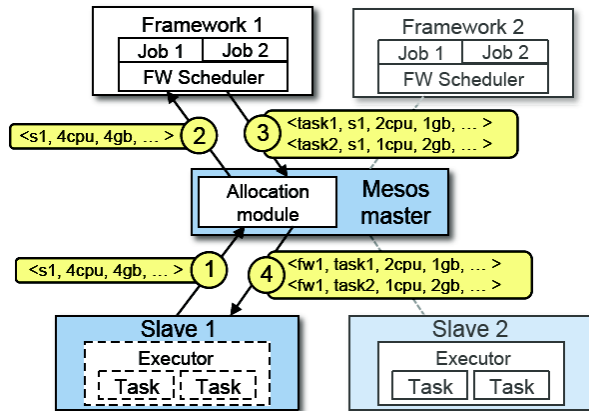
- **Slaves** continuously send status updates about **resources** to the **Master**.

## Mesos Architecture (2/4)




- Pluggable **scheduler** picks **framework** to send an **offer** to.

## Mesos Architecture (3/4)



- Framework scheduler selects resources and provides tasks.

## Mesos Architecture (4/4)




- Framework executors launch tasks.

## Question?


How to allocate resources of different types?

# Single Resource: Fair Sharing


- ▶  $n$  users want to share a resource, e.g., CPU.
  - **Solution:** allocate each  $\frac{1}{n}$  of the shared resource.



- ▶ Generalized by **max-min fairness**.
  - Handles if a user wants less than its fair share.
  - E.g., user 1 wants no more than 20%.



- ▶ Generalized by **weighted max-min fairness**.
  - Give **weights** to users according to **importance**.
  - E.g., user 1 gets weight 1, user 2 weight 2.







## Max-Min Fairness - Example

- ▶ 1 resource: CPU
- ▶ Total resources: 20 CPU
- ▶ User 1 has  $x$  tasks and wants  $\langle 1\text{CPU} \rangle$  per task
- ▶ User 2 has  $y$  tasks and wants  $\langle 2\text{CPU} \rangle$  per task

$\max(x, y)$  (maximize allocation)

subject to

$x + 2y \leq 20$  (CPU constraint)

$x = 2y$

so

$x = 10$

$y = 5$

# Properties of Max-Min Fairness

## ► Share guarantee

- Each user can get at least  $\frac{1}{n}$  of the resource.
- But will get less if her demand is less.

## ► Strategy proof

- Users are not better off by asking for more than they need.
- Users have no reason to lie.

► Max-Min fairness is the only reasonable mechanism with these two properties.


► Widely used: OS, networking, datacenters, ...

## Question?

When is Max-Min Fairness **NOT** Enough?

Need to schedule **multiple, heterogeneous** resources, e.g.,  
CPU, memory, etc.

- 1 resource: CPU
- User 1 wants  $\langle 1\text{CPU} \rangle$  per task
- User 2 wants  $\langle 2\text{CPU} \rangle$  per task



- 2 resources: CPUs and mem
- User 1 wants  $\langle 1\text{CPU}, 4\text{GB} \rangle$  per task
- User 2 wants  $\langle 2\text{CPU}, 1\text{GB} \rangle$  per task
- What is a fair allocation?

## A Natural Policy (1/2)

- ▶ **Asset fairness**: give weights to resources (e.g., 1 CPU = 1 GB) and **equalize total value given to each user**.
- ▶ Total resources: 28 CPU and 56GB RAM (e.g., 1 CPU = 2 GB)
  - User 1 has  $x$  tasks and wants  $\langle 1\text{CPU}, 2\text{GB} \rangle$  per task
  - User 2 has  $y$  tasks and wants  $\langle 1\text{CPU}, 4\text{GB} \rangle$  per task
- ▶ Asset fairness yields:

$$\max(x, y)$$


$$x + y \leq 28$$

$$2x + 4y \leq 56$$


$$2x = 3y$$

User 1:  $x = 12$ :  $\langle 43\%\text{CPU}, 43\%\text{GB} \rangle$  ( $\sum = 86\%$ )

User 2:  $y = 8$ :  $\langle 28\%\text{CPU}, 57\%\text{GB} \rangle$  ( $\sum = 86\%$ )



## A Natural Policy (2/2)



- ▶ **Problem:** violates share grantee.
- ▶ User 1 gets less than 50% of both CPU and RAM.
- ▶ Better off in a separate cluster with half the resources.



# Challenge

- ▶ Can we find a fair sharing policy that provides:
  - Share guarantee
  - Strategy-proofness
  
- ▶ Can we generalize max-min fairness to multiple resources?



## Proposed Solution

Dominant Resource Fairness (**DRF**)



## Dominant Resource Fairness (DRF) (1/2)

- ▶ **Dominant resource** of a user: the resource that user has the **biggest share of**.
  - Total resources:  $\langle 8\text{CPU}, 5\text{GB} \rangle$
  - User 1 allocation:  $\langle 2\text{CPU}, 1\text{GB} \rangle$ :  $\frac{2}{8} = 25\%$  CPU and  $\frac{1}{5} = 20\%$  RAM
  - Dominant resource of User 1 is **CPU** ( $25\% > 20\%$ )
- ▶ **Dominant share** of a user: the **fraction** of the **dominant resource** she is allocated.
  - User 1 dominant share is **25%**.

## Dominant Resource Fairness (DRF) (2/2)

- ▶ Apply **max-min fairness** to **dominant shares**: give every user an equal share of her dominant resource.
- ▶ **Equalize** the **dominant share** of the users.
  - Total resources:  $\langle 9\text{CPU}, 18\text{GB} \rangle$
  - User 1 wants  $\langle 1\text{CPU}, 4\text{GB} \rangle$ ; Dominant resource: RAM ( $\frac{1}{9} < \frac{4}{18}$ )
  - User 2 wants  $\langle 3\text{CPU}, 1\text{GB} \rangle$ ; Dominant resource: CPU ( $\frac{3}{9} > \frac{1}{18}$ )
- ▶  $\max(x, y)$ 


$$x + 3y \leq 9$$

$$4x + y \leq 18$$

$$\frac{4x}{18} = \frac{3y}{9}$$

User 1:  $x = 3$ :  $\langle 33\%\text{CPU}, 66\%\text{GB} \rangle$

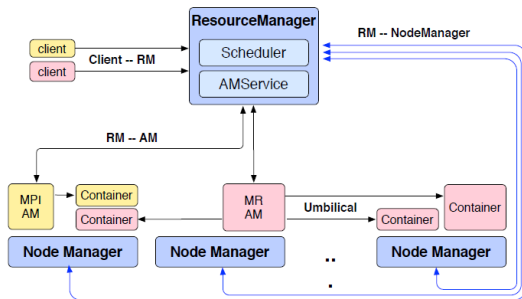
User 2:  $y = 2$ :  $\langle 66\%\text{CPU}, 16\%\text{GB} \rangle$



# YARN


# YARN Architecture

- ▶ Resource Manager (RM)
- ▶ Application Master (AM)
- ▶ Node Manager (NM)




# YARN Architecture - Resource Manager (1/2)

- ▶ One per cluster
  - Central: global view
- ▶ Job requests are submitted to RM.
  - To start a job, RM finds a container to spawn AM.
- ▶ Container: logical bundle of resources (CPU/memory)



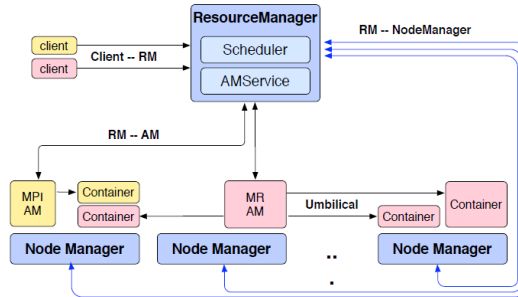
## YARN Architecture - Resource Manager (2/2)

- ▶ Only handles an **overall resource** profile for **each job**.
  - **Local optimization** is up to the job.
- ▶ Preemption
  - **Request resources back** from an job.
  - **Checkpoint** jobs




# YARN Architecture - Application Manager

- ▶ The head of a job.
- ▶ Runs as a container.
- ▶ Request resources from RM (num. of containers/resource per container/locality ...)



# YARN Architecture - Node Manager (1/2)


- ▶ The **worker daemon**.
- ▶ Registers with RM.
- ▶ **One** per node.
- ▶ **Report resources** to RM: memory, CPU, ...





## YARN Architecture - Node Manager (2/2)

- ▶ Configure the environment for task execution.
- ▶ Garbage collection.
- ▶ Auxiliary services.
  - A process may produce data that persist beyond the life of the container.
  - Output intermediate data between map and reduce tasks.



## YARN Framework (1/2)

- ▶ **Containers** are described by a **Container Launch Context (CLC)**.
  - The command necessary to create the process
  - Environment variables
  - Security tokens
  - ...
- ▶ **Submitting the job**: passing a **CLC** for the **AM** to the **RM**.
- ▶ When **RM** starts the **AM**, it should register with the **RM**.
  - Periodically advertise its **liveness** and **requirements** over the **heartbeat** protocol.



## YARN Framework (2/2)

- ▶ Once the **RM** allocates a container, **AM** can construct a **CLC** to launch the container on the corresponding **NM**.
  - It **monitors** the status of the **running container** and stop it when the resource should be reclaimed.
- ▶ Once the **AM** is done with its work, it should unregister from the **RM** and **exit cleanly**.

# Summary



# Summary

## ▶ Mesos

- Offered-based
- Max-Min fairness: DRF

## ▶ YARN

- Request-based
- RM, AM, NM



## References

- ▶ B. Hindman et al., “Mesos: A Platform for Fine-Grained Resource Sharing in the Data Center”, NSDI 2011
- ▶ V. Vavilapalli et al., “Apache hadoop yarn: Yet another resource negotiator”, ACM Cloud Computing 2013

# Questions?

## Acknowledgements

Some slides were derived from Ion Stoica and Ali Ghodsi slides (Berkeley University), and Wei-Chiu Chuang slides (Purdue University).