



Introduction to Data Stream Processing

Amir H. Payberah
payberah@kth.se
2021-09-21



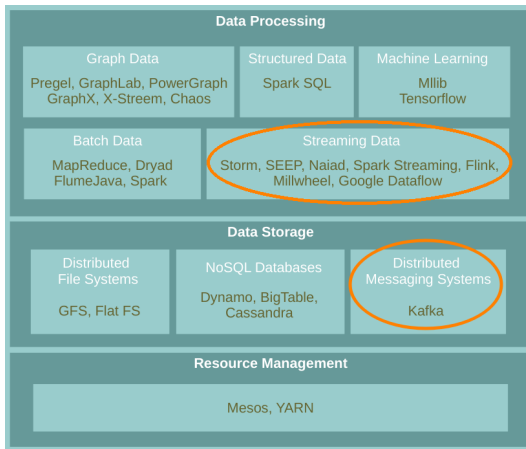


The Course Web Page

<https://id2221kth.github.io>

<https://tinyurl.com/f6x544h>

Where Are We?




Stream Processing (1/4)

- **Stream processing** is the act of **continuously** incorporating **new data** to compute a result.




Stream Processing (2/4)

- ▶ The **input data** is **unbounded**.
 - A **series of events**, no predetermined **beginning or end**.
 - E.g., credit card transactions, clicks on a website, or sensor readings from IoT devices.



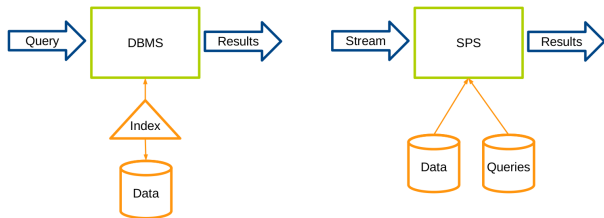
Stream Processing (3/4)

- ▶ **User applications** can then compute **various queries** over this stream of events.
 - E.g., **tracking** a running count of each type of event, or **aggregating** them into hourly windows.

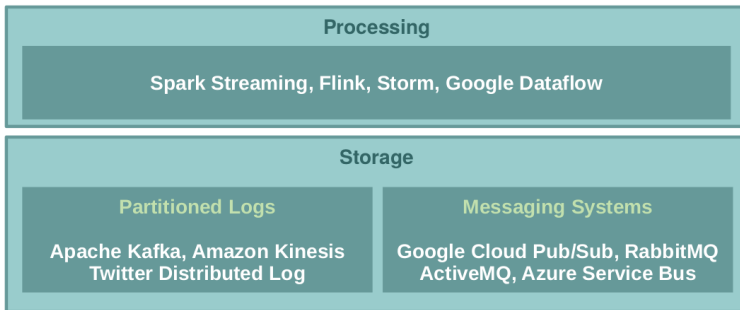


Stream Processing (4/4)

- ▶ Database Management Systems (DBMS): **data-at-rest** analytics
 - **Store** and **index** data before processing it.
 - Process data only when **explicitly** asked by the users.
- ▶ Stream Processing Systems (SPS): **data-in-motion** analytics
 - Processing information as it **flows**, **without storing** them persistently.



Stream Processing Systems Stack

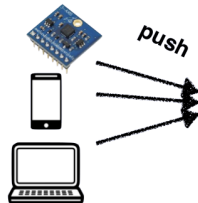


Data Stream Storage


The Problem

- ▶ We need disseminate streams of events from various producers to various consumers.

Data Producers



Data Consumers






Example

- ▶ Suppose you have a [website](#), and every time someone [loads a page](#), you send a [viewed page](#) event to consumers.
- ▶ The consumers may do any of the following:
 - [Store](#) the message in HDFS for future analysis
 - [Count page](#) views and update a dashboard
 - Trigger an [alert](#) if a page view fails
 - Send an [email](#) notification to another user

Possible Solution?

- ▶ Messaging systems



Message

www.defit.org




What is Messaging System?

- ▶ **Messaging system** is an approach to **notify consumers** about new events.
- ▶ **Messaging systems**
 - **Direct** messaging
 - Message **brokers**


Direct Messaging (1/2)

- ▶ Necessary in **latency critical** applications (e.g., remote surgery).
- ▶ A **producer** sends a message containing the event, which is **pushed** to **consumers**.
- ▶ Both consumers and producers have to be **online at the same time**.




Direct Messaging (2/2)

- ▶ What happens if a **consumer crashes** or temporarily **goes offline**? (**not durable**)
- ▶ What happens if **producers** send messages **faster** than the **consumers** can process?
 - **Dropping** messages
 - **Backpressure**
- ▶ We need **message brokers** that can **log events** to process at a **later time**.




Message Broker



[<https://bluesyemre.com/2018/10/16/thousands-of-scientists-publish-a-paper-every-five-days>]


Message Broker

- ▶ A **message broker** decouples the **producer-consumer** interaction.
- ▶ It runs as a **server**, with **producers and consumers** connecting to it as **clients**.
- ▶ **Producers** write messages to the broker, and **consumers** receive them by reading them from the broker.
- ▶ **Consumers** are generally **asynchronous**.




Message Broker (2/2)

- ▶ When **multiple consumers** read messages in the **same topic**.
- ▶ **Load balancing**: each message is delivered to **one** of the consumers.



- ▶ **Fan-out**: each message is delivered to **all** of the consumers.






Partitioned Logs (1/2)

- ▶ In typical message brokers, once a message is **consumed**, it is **deleted**.
- ▶ **Log-based message brokers** **durably** store all events in a sequential **log**.
- ▶ A **log** is an **append-only** sequence of records on **disk**.
- ▶ A **producer** sends a message by **appending** it to the end of the log.
- ▶ A **consumer** receives messages by reading the log **sequentially**.

Partitioned Logs (2/2)

- ▶ To **scale up** the system, logs can be **partitioned** hosted on **different machines**.
- ▶ Each **partition** can be read and written **independently** of others.
- ▶ A **topic** is a **group of partitions** that all carry messages of the **same type**.
- ▶ **Within each partition**, the broker assigns a **monotonically increasing sequence number (offset)** to every message
- ▶ **No ordering** guarantee **across partitions**.




Kafka - A Log-Based Message Broker




Kafka (1/5)

- **Kafka** is a distributed, topic oriented, partitioned, replicated commit **log service**.




Kafka (2/5)

- Kafka is a **distributed**, topic oriented, partitioned, replicated commit **log service**.




Kafka (3/5)

- Kafka is a distributed, topic oriented, partitioned, replicated commit log service.




Kafka (4/5)

- Kafka is a distributed, topic oriented, partitioned, replicated commit log service.



Kafka (5/5)

- Kafka is a distributed, topic oriented, partitioned, replicated commit log service.



Logs, Topics and Partition (1/5)


- Kafka is about **logs**.
- **Topics** are **queues**: a **stream of messages** of a **particular type**

```
jkreps-mn:~ jkreps$ tail -f -n 20 /var/log/apache2/access_log
::1 - - [23/Mar/2014:15:07:00 -0700] "GET /images/apache_feather.gif HTTP/1.1" 200 4128
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/producer_consumer.png HTTP/1.1" 200 8f
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/log_anatomy.png HTTP/1.1" 200 19579
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/consumer-groups.png HTTP/1.1" 200 268;
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/log_compaction.png HTTP/1.1" 200 4141;
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /documentation.html HTTP/1.1" 200 189893
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/log_cleaner_anatomy.png HTTP/1.1" 200
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/kafka_log.png HTTP/1.1" 200 134321
::1 - - [23/Mar/2014:15:07:04 -0700] "GET /images/mirror-maker.png HTTP/1.1" 200 17054
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /documentation.html HTTP/1.1" 200 189937
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /styles.css HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/kafka_logo.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/producer_consumer.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/log_anatomy.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/consumer-groups.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/log_cleaner_anatomy.png HTTP/1.1" 304
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/log_compaction.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/kafka_log.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:08:07 -0700] "GET /images/mirror-maker.png HTTP/1.1" 304 -
::1 - - [23/Mar/2014:15:09:55 -0700] "GET /documentation.html HTTP/1.1" 200 195264
```

0	1	2	3	4	5	6	7	8	9	10	11	12
---	---	---	---	---	---	---	---	---	---	----	----	----

Logs, Topics and Partition (2/5)


- Each **message** is assigned a **sequential id** called an **offset**.



Logs, Topics and Partition (3/5)


► **Topics** are **logical** collections of **partitions** (the **physical files**).

- Ordered
- Append only
- Immutable




Logs, Topics and Partition (4/5)

- ▶ Ordering is only **guaranteed within** a **partition** for a **topic**.
- ▶ Messages sent by a **producer** to a particular topic partition will be **appended** in the order they are sent.
- ▶ A **consumer** instance sees messages in the order they are stored in the log.




Logs, Topics and Partition (5/5)

- ▶ **Partitions** of a topic are **replicated**: **fault-tolerance**
- ▶ A **broker** contains some of the **partitions** for a topic.
- ▶ One broker is the **leader** of a partition: all **writes** and **reads** must go to the leader.



Kafka Architecture



Coordination

- ▶ Kafka uses **Zookeeper** for the following tasks:
- ▶ Detecting the **addition** and the **removal** of **brokers** and **consumers**.
- ▶ Keeping track of the **consumed** offset of each partition.





State in Kafka

- ▶ Brokers are **stateless**: **no metadata** for consumers-producers in **brokers**.
- ▶ **Consumers** are responsible for keeping track of **offsets**.
- ▶ Messages in queues **expire** based on pre-configured time periods (e.g., once a day).



Delivery Guarantees

- ▶ Kafka guarantees that messages from a **single partition** are delivered to a consumer **in order**.
- ▶ There is **no guarantee** on the ordering of messages coming from **different partitions**.
- ▶ Kafka only guarantees **at-least-once** delivery.



Start and Work With Kafka

```
# Start the ZooKeeper
```

```
zookeeper-server-start.sh config/zookeeper.properties
```

```
# Start the Kafka server
```

```
kafka-server-start.sh config/server.properties
```

```
# Create a topic, called "avg"
```

```
kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1  
--topic avg
```

```
# Produce messages and send them to the topic "avg"
```

```
kafka-console-producer.sh --broker-list localhost:9092 --topic avg
```

```
# Consume the messages sent to the topic "avg"
```

```
kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic avg --from-beginning
```

Data Stream Processing



Streaming Data

- ▶ Data stream is unbound data, which is broken into a sequence of individual tuples.
- ▶ A data tuple is the atomic data item in a data stream.
- ▶ Can be structured, semi-structured, and unstructured.



Streaming Data Processing Design Points

- ▶ Continuous vs. micro-batch processing
- ▶ Record-at-a-Time vs. declarative APIs
- ▶ Event time vs. processing time
- ▶ Windowing



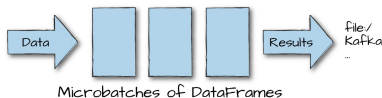
Streaming Data Processing Design Points

- ▶ Continuous vs. micro-batch processing
- ▶ Record-at-a-Time vs. declarative APIs
- ▶ Event time vs. processing time
- ▶ Windowing

Streaming Data Processing Patterns

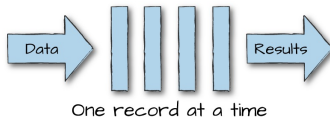
► Micro-batch systems

- Batch engines
- Slicing up the unbounded data into a sets of bounded data, then process each batch.



► Continuous processing-based systems

- Each node in the system continually listens to messages from other nodes and outputs new updates to its child nodes.





Streaming Data Processing Design Points

- ▶ Continuous vs. micro-batch processing
- ▶ Record-at-a-Time vs. declarative APIs
- ▶ Event time vs. processing time
- ▶ Windowing

Record-at-a-Time vs. Declarative APIs

► Record-at-a-Time API (e.g., Storm)

- Low-level API
- Passes **each event** to the **application** and let it react.
- Useful when applications need **full control** over the processing of data.
- **Complicated factors**, such as maintaining state, are **governed by the application**.

► Declarative API (e.g., Spark streaming, Flink, Google Dataflow)

- Applications specify **what** to compute **not how** to compute it in response to **each new event**.



Streaming Data Processing Design Points

- ▶ Continuous vs. micro-batch processing
- ▶ Record-at-a-Time vs. declarative APIs
- ▶ Event time vs. processing time
- ▶ Windowing




Event Time vs. Processing Time (1/2)

- ▶ **Event time**: the time at which events **actually occurred**.
 - Timestamps inserted into each record **at the source**.
- ▶ **Processing time**: the time when the record is **received at the streaming application**.

Event Time vs. Processing Time (2/2)

- ▶ Ideally, event time and processing time should be equal.
- ▶ Skew between event time and processing time.



[<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>]



Streaming Data Processing Design Points

- ▶ Continuous vs. micro-batch processing
- ▶ Record-at-a-Time vs. declarative APIs
- ▶ Event time vs. processing time
- ▶ Windowing



Windowing (1/2)


- ▶ **Window**: a **buffer** associated with an input port to retain previously **received tuples**.
- ▶ **Four** different windowing **management policies**.
 - **Count-based policy**: the **maximum number** of tuples a window buffer can hold
 - **Delta-based policy**: a **delta threshold** in a tuple attribute
 - **Punctuation-based policy**: a **punctuation** is received
 - **Time-based policy**: based on **processing or event time** period

Windowing (2/2)

- ▶ Two types of windows: **tumbling** and **sliding**
- ▶ **Tumbling window**: supports **batch** operations.
 - When the buffer fills up, **all** the tuples are **evicted**.

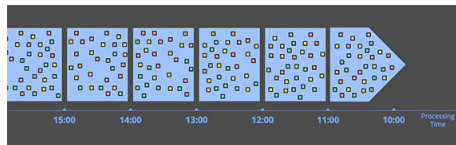


- ▶ **Sliding window**: supports **incremental** operations.
 - When the buffer fills up, **older** tuples are **evicted**.



Windowing by Processing Time


- ▶ The system **buffers up** incoming data into windows until **some amount of processing time has passed**.
- ▶ E.g., **five-minute** fixed windows



[<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>]


Windowing by Event Time

- ▶ Reflect the **times** at which **events** actually happened.
- ▶ Handling **out-of-order** events.




[<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>]

- ▶ **Watermarking** helps a stream processing system to deal with **lateness**.
- ▶ Watermarks **flow as part of the data stream** and carry a **timestamp t**.
- ▶ A watermark is a **threshold** to specify **how long the system waits** for **late events**.
- ▶ Streaming systems uses **watermarks** to **measure progress** in **event time**.



Windowing by Event Time - Watermark (2/2)


- ▶ A $W(t)$ declares that **event time** has reached time t in that stream
 - There should be **no more elements from the stream** with a timestamp $t' \leq t$.
- ▶ It is possible that certain elements will **violate the watermark condition**.
 - After the $W(t)$ has occurred, more elements with timestamp $t' \leq t$ will occur.
- ▶ If an arriving event lies **within the watermark**, it gets used to update a query.
- ▶ Streaming programs may explicitly expect some **late elements**.



Streaming Data Processing Model

Streaming Data Processing

- ▶ The tuples are processed by the application's **operators** or **processing element (PE)**.
- ▶ A **PE** is the **basic functional unit** in an application.
 - A PE processes **input** tuples, applies a **function**, and **outputs** tuples.
 - A **set of PEs** and stream **connections**, organized into a **data flow graph**.





PEs States (1/3)

- ▶ A PE can either **maintain internal state** across tuples while processing them, or process tuples **independently** of each other.
- ▶ **Stateful** vs. **stateless** tasks



PEs States (2/3)

- ▶ **Stateless** tasks: do **not maintain state** and process each tuple **independently** of **prior history**, or even from the **order** of arrival of tuples.
- ▶ Easily **parallelized**.
- ▶ **No synchronization**.
- ▶ **Restart upon failures** without the need of any recovery procedure.



PEs States (3/3)

- ▶ **Stateful** tasks: involves **maintaining** information **across different tuples** to detect complex patterns.
- ▶ A **PE** is usually a **synopsis** of the **tuples** received so far.
- ▶ A subset of **recent tuples** kept in a **window buffer**.


Runtime Systems




Job and Job Management

- ▶ At runtime, an **application** is represented by **one or more jobs**.
- ▶ **Jobs** are deployed as a **collection of PEs**.
- ▶ **Job management** component must **identify and track** individual **PEs**, the **jobs** they belong to, and associate them with the user that instantiated them.



-




Logical Plan vs. Physical Plan (2/3)



Logical plan



Different physical plans





Logical Plan vs. Physical Plan (3/3)

- ▶ How to map a **network of PEs** onto the **physical network of nodes**?
 - Parallelization
 - Fault tolerance
 - Optimization

Parallelization




Parallelization

- ▶ How to **scale** with increasing the **number queries** and the **rate of incoming events**?
- ▶ **Three** forms of parallelisms.
 - **Pipelined** parallelism
 - **Task** parallelism
 - **Data** parallelism


Pipelined Parallelism


- Sequential stages of a computation execute **concurrently** for **different data items**.



Task Parallelism


- Independent processing stages of a larger computation are executed **concurrently** on the same or distinct data items.





Data Parallelism (2/2)

- How to **allocate data items** to each **computation instance**?



Fault Tolerance



Fault Tolerance

- ▶ The recovery methods of streaming frameworks must take:
 - **Correctness**, e.g., data loss and duplicates
 - **Performance**, e.g., low latency



Delivery Guarantees

- ▶ **At-least-once**: might appear **many times**
- ▶ **Exactly-once**: is consumed just **once**




Recovery Methods

- ▶ Active backup
- ▶ Passive backup
- ▶ Upstream backup

Recovery Methods - Active Backup

- ▶ Each processing node has an associated **backup node**.
- ▶ Both primary and backup nodes are given the **same** input.
- ▶ If the **primary** fails, the **backup** takes over by **sending the logged tuples** to all downstream neighbors and then continuing its processing.






Recovery Methods - Passive Backup

- ▶ Periodically check-points processing state to a shared storage.
- ▶ The backup node takes over from the latest checkpoint when the primary fails.

Recovery Methods - Upstream Backup

- ▶ Upstream nodes store the tuples until the downstream nodes acknowledge them.
- ▶ If a node fails, an empty node rebuilds the latest state of the failed primary from the logs kept at the upstream server.
- ▶ There is no backup node in this model.



Summary



Summary

- ▶ Messaging system and partitioned logs
- ▶ Decoupling producers and consumers
- ▶ Kafka: distributed, topic oriented, partitioned, replicated log service
- ▶ Logs, topics, partition
- ▶ Kafka architecture: producer, consumer, broker, coordinator



Summary

- ▶ SPS vs. DBMS
- ▶ Data stream, unbounded data, tuples
- ▶ Event-time vs. processing time
- ▶ Micro-batch vs. continuous processing (windowing)
- ▶ PEs and dataflow
- ▶ Stateless vs. Stateful PEs
- ▶ SPS runtime: parallelization, fault-tolerance

- ▶ J. Kreps et al., “Kafka: A distributed messaging system for log processing”, NetDB 2011
- ▶ M. Zaharia et al., “Spark: The Definitive Guide”, O'Reilly Media, 2018 - Chapter 20
- ▶ M. Fraggkoulis et al., “A Survey on the Evolution of Stream Processing Systems”, 2020
- ▶ J. Hwang et al., “High-availability algorithms for distributed stream processing”, ICDE 2005
- ▶ T. Akidau, “The world beyond batch: Streaming 101”,
<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>

Questions?