



Scalable Stream Processing - Spark Streaming

Amir H. Payberah
payberah@kth.se
2023-09-25





The Course Web Page

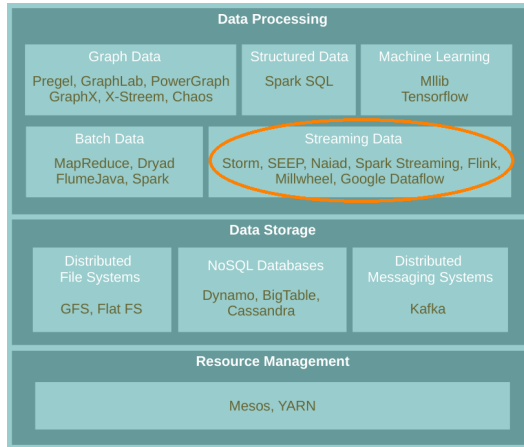
<https://id2221kth.github.io>



The Questions-Answers Page

<https://tinyurl.com/hk7hzpw5>

Where Are We?



Spark Streaming


Spark Streaming

- ▶ Run a streaming computation as a **series** of very **small**, **deterministic batch jobs**.
 - **Chops up** the live stream into batches of **X** seconds.
 - Treats each batch as **RDDs** and processes them using **RDD operations**.
 - Discretized Stream Processing (**DStream**)




DStream (1/2)

- **DStream**: sequence of **RDDs** representing a stream of data.



DStream (2/2)

- Any **operation** applied on a **DStream** translates to operations on the underlying **RDDs**.





StreamingContext

- ▶ **StreamingContext** is the **main entry** point of all Spark Streaming functionality.

```
val conf = new SparkConf().setAppName(appName).setMaster(master)
val ssc = new StreamingContext(conf, Seconds(1))
```

- ▶ The second parameter, **Seconds(1)**, represents the **time interval** at which streaming data will be divided into **batches**.



Input Operations

► Socket connection

- Creates a DStream from text data received over a **TCP socket connection**.

```
ssc.socketTextStream("localhost", 9999)
```

► File stream

- Reads data from **files**.

```
streamingContext.fileStream[KeyClass, ValueClass, InputFormatClass](dataDirectory)  
streamingContext.textFileStream(dataDirectory)
```

► Connectors with external sources, e.g., **Twitter, Kafka, Flume, Kinesis, ...**



Transformations (1/2)

- ▶ Transformations on DStreams are still lazy!
- ▶ DStreams support many of the transformations available on normal Spark RDDs.
- ▶ Computation is kicked off explicitly by a call to the `start()` method.



Transformations (2/2)

- ▶ **map**: a new **DStream** by passing each **element** of the source DStream through a given function.
- ▶ **reduce**: a new DStream of **single-element RDDs** by **aggregating** the elements in each RDD using a given function.
- ▶ **reduceByKey**: a new DStream of **(K, V) pairs** where the values for each key are **aggregated** using the given reduce function.



Example - Word Count (1/6)

- First we create a `StreamingContext`

```
import org.apache.spark._  
import org.apache.spark.streaming._  
  
// Create a local StreamingContext with two working threads and batch interval of 1 second.  
val conf = new SparkConf().setMaster("local[2]").setAppName("NetworkWordCount")  
val ssc = new StreamingContext(conf, Seconds(1))
```



Example - Word Count (2/6)


- ▶ Create a **DStream** that represents streaming data from a **TCP source**.
- ▶ Specified as **hostname** (e.g., localhost) and **port** (e.g., 9999).

```
val lines = ssc.socketTextStream("localhost", 9999)
```

Example - Word Count (3/6)

- ▶ Use `flatMap` on the stream to split the records text to words.
- ▶ It creates a new DStream.

```
val words = lines.flatMap(_.split(" "))
```





Example - Word Count (4/6)

- ▶ Map the **words** DStream to a DStream of **(word, 1)**.
- ▶ Get the **frequency of words** in each **batch of data**.
- ▶ Finally, **print** the result.

```
val pairs = words.map(word => (word, 1))  
val wordCounts = pairs.reduceByKey(_ + _)  
wordCounts.print()
```




Example - Word Count (5/6)

- Start the **computation** and **wait** for it to **terminate**.


```
// Start the computation  
ssc.start()  
  
// Wait for the computation to terminate  
ssc.awaitTermination()
```

Example - Word Count (6/6)

```
val conf = new SparkConf().setMaster("local[2]").setAppName("NetworkWordCount")
val ssc = new StreamingContext(conf, Seconds(1))


val lines = ssc.socketTextStream("localhost", 9999)
val words = lines.flatMap(_.split(" "))
val pairs = words.map(word => (word, 1))
val wordCounts = pairs.reduceByKey(_ + _)
wordCounts.print()

ssc.start()
ssc.awaitTermination()
```



Window Operations (1/2)

- ▶ Spark provides a set of transformations that apply to a over a **sliding window** of data.
- ▶ A window is defined by two parameters: **window length** and **slide interval**.
- ▶ A **tumbling window** effect can be achieved by making **slide interval = window length**



Window Operations (2/2)

- ▶ `reduceByWindow(func, windowLength, slideInterval)`
 - Returns a new **single-element DStream**, created by aggregating elements in the stream over a **sliding interval** using `func`.

- ▶ `reduceByKeyAndWindow(func, windowLength, slideInterval)`
 - Called on a DStream of **(K, V) pairs**.
 - Returns a **new DStream of (K, V) pairs** where the values for each key are aggregated using function `func` over **batches in a sliding window**.

Example - Word Count with Window

```
val conf = new SparkConf().setMaster("local[2]").setAppName("NetworkWordCount")
val ssc = new StreamingContext(conf, Seconds(1))

val lines = ssc.socketTextStream("localhost", 9999)
val words = lines.flatMap(_.split(" "))
val pairs = words.map(word => (word, 1))
val windowedWordCounts = pairs.reduceByKeyAndWindow(_ + _, Seconds(30), Seconds(10))
windowedWordCounts.print()

ssc.start()
ssc.awaitTermination()
```





What about States?

- ▶ Accumulate and aggregate the results from the start of the streaming job.
- ▶ Need to check the previous state of the RDD in order to do something with the current RDD.
- ▶ Spark supports stateful streams.



Checkpointing

- It is **mandatory** that you provide a checkpointing directory for **stateful streams**.

```
val ssc = new StreamingContext(conf, Seconds(1))  
ssc.checkpoint("path/to/persistent/storage")
```



Stateful Stream Operations

► `mapWithState`

- It is executed only on set of keys that are available in the `last micro batch`.

```
def mapWithState[StateType, MappedType](spec: StateSpec[K, V, StateType, MappedType]):  
    DStream[MappedType]  
  
StateSpec.function(updateFunc)  
val updateFunc = (batch: Time, key: String, value: Option[Int], state: State[Int])
```

- ## ► Define the update function (`partial updates`) in `StateSpec`.

Example - Stateful Word Count (1/4)

```
val ssc = new StreamingContext(conf, Seconds(1))
ssc.checkpoint(".")

val lines = ssc.socketTextStream(IP, Port)
val words = lines.flatMap(_.split(" "))
val pairs = words.map(word => (word, 1))

val stateWordCount = pairs.mapWithState(StateSpec.function(updateFunc))

val updateFunc = (key: String, value: Option[Int], state: State[Int]) => {
  val newCount = value.getOrElse(0)
  val oldCount = state.getOption.getOrElse(0)
  val sum = newCount + oldCount
  state.update(sum)
  (key, sum)
}
```



Example - Stateful Word Count (2/4)

- ▶ The first micro batch contains a message `a`.
- ▶ `updateFunc = (key: String, value: Option[Int], state: State[Int]) => (key, sum)`
- ▶ Input: `key = a, value = Some(1), state = 0`
- ▶ Output: `key = a, sum = 1`

Example - Stateful Word Count (3/4)

- ▶ The **second micro batch** contains messages **a** and **b**.
- ▶ `updateFunc = (key: String, value: Option[Int], state: State[Int]) => (key, sum)`
- ▶ Input: `key = a, value = Some(1), state = 1`
- ▶ Input: `key = b, value = Some(1), state = 0`
- ▶ Output: `key = a, sum = 2`
- ▶ Output: `key = b, sum = 1`



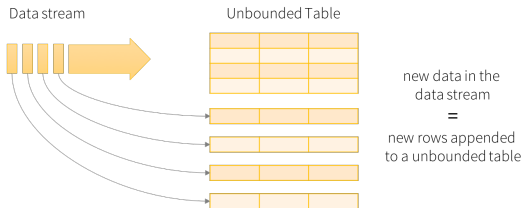
Example - Stateful Word Count (4/4)

- ▶ The **third micro batch** contains a message **b**.
- ▶ `updateFunc = (key: String, value: Option[Int], state: State[Int]) => (key, sum)`
- ▶ Input: `key = b, value = Some(1), state = 1`
- ▶ Output: `key = b, sum = 2`

Structured Streaming

Structured Streaming

- Treating a **live data stream** as a **table** that is being **continuously appended**.




Data stream as an unbounded table




Programming Model (1/2)

- ▶ Defines a **query** on the input table, as a **static table**.
 - Spark automatically converts this **batch-like query** to a **streaming execution plan**.
- ▶ Specify **triggers** to control **when to update the results**.
 - Each time a trigger fires, Spark **checks for new data** (**new row** in the input table), and **incrementally** updates the result.

Programming Model (2/2)



User's batch-like
query on input table



Incremental execution on streaming data

► Three output modes:

1. **Append**: only the new rows **appended to the result table** since the last trigger will be written to the external storage.
2. **Complete**: the **entire updated result table** will be written to external storage.
3. **Update**: only the rows that were **updated in the result table** since the last trigger will be changed in the external storage.



Steps to Define a Streaming Query (1/4)

- ▶ Define **input sources**.
- ▶ Use `spark.readStream` to create a `DataStreamReader`.

```
val spark = SparkSession.builder.master("local[2]").appName("appname").getOrCreate()

val lines = spark.readStream.format("socket")
    .option("host", "localhost")
    .option("port", 9999)
    .load()
```



Steps to Define a Streaming Query (2/4)

- ▶ Transform data.
- ▶ E.g., below `counts` is a `streaming DataFrame` that represents the running word counts.

```
import org.apache.spark.sql.functions._  
  
val words = lines.select(split(col("value"), " ").as("word"))  
  
val wordCounts = words.groupBy("word").count()
```

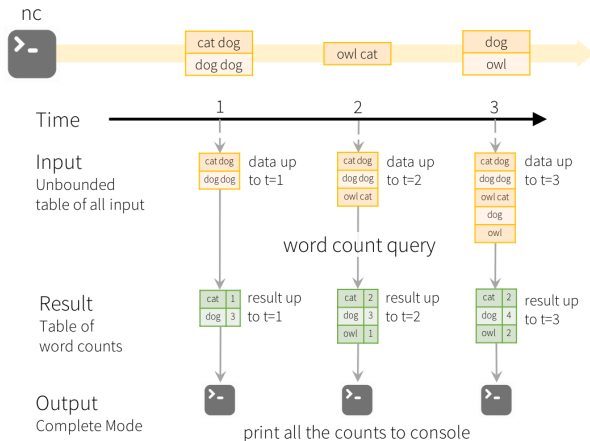


Steps to Define a Streaming Query (3/4)

- ▶ Define **output sink** and **output mode**.
- ▶ Use `DataFrame.writeStream` to define how to write the processed output data.
- ▶ **Start** the query.

```
val query = wordCounts.writeStream.format("console").outputMode("complete").start()  
query.awaitTermination()
```

Steps to Define a Streaming Query (4/4)



[<https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>]

Streaming Data Sources and Sinks - Files (1/2)

► Reading from files.

```
import org.apache.spark.sql.types._

val inputDirectoryOfJsonFiles = ...

val fileSchema = new StructType()
  .add("key", IntegerType)
  .add("value", IntegerType)

val inputDF = spark.readStream
  .format("json")
  .schema(fileSchema)
  .load(inputDirectoryOfJsonFiles)
```

Streaming Data Sources and Sinks - Files (2/2)

► Writing to files.

```
val outputDir = ...  
val checkpointDir = ...  
val resultDF = ...  
  
val streamingQuery = resultDF  
  .writeStream  
  .format("parquet")  
  .option("path", outputDir)  
  .option("checkpointLocation", checkpointDir)  
  .start()
```



Basic Operations (1/2)

- Most of operations on `DataFrame/Dataset` are **supported** for streaming.

```
case class Call(action: String, time: Timestamp, id: Int)

val df: DataFrame = spark.readStream.json("s3://logs")
val ds: Dataset[Call] = df.as[Call]
```

- Selection and projection

```
df.select("action").where("id > 10") // using untyped APIs
ds.filter(_.id > 10).map(_.action) // using typed APIs
```




Basic Operations (2/2)


► Aggregation

```
df.groupBy("action") // using untyped API  
ds.groupByKey(_.action) // using typed API
```

► SQL commands

```
df.createOrReplaceTempView("dfView")  
spark.sql("select count(*) from dfView") // returns another streaming DF
```


- Computing counts corresponding to 10-minute windows sliding every five minutes.



```
words = ... // streaming DataFrame of schema { timestamp: Timestamp, word: String }

windowedCounts = words.groupBy( \
  window($"timestamp", "10 minutes", "5 minutes"), $"word").count()
```

- The **watermark delay** defines **how long** the engine will **wait** for **late data** to arrive.



```
words = ... # streaming DataFrame of schema { timestamp: Timestamp, word: String }

windowedCounts = words.withWatermark("timestamp", "10 minutes") \
    .groupBy( \
        window($"timestamp", "10 minutes", "5 minutes"), $"word").count()
```



Summary

- ▶ Mini-batch processing
- ▶ DStream: sequence of RDDs
- ▶ RDD and window operations
- ▶ Structured streaming

- ▶ M. Zaharia et al., “Spark: The Definitive Guide”, O'Reilly Media, 2018 - Chapters 20-23.
- ▶ M. Zaharia et al., “Discretized Streams: An Efficient and Fault-Tolerant Model for Stream Processing on Large Clusters”, HotCloud'12.
- ▶ The world beyond batch: Streaming 102
<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-102>

Questions?